

Parallel Graph Coloring

Andrea Cavallo, Gioele Scaletta

System and Device Programming
Politecnico di Torino

September 3rd, 2021

Table of Contents

- ① Introduction
- ② Data Structure
- ③ Algorithms
- ④ Parallelization strategy
- ⑤ Synchronization strategy
- ⑥ Hyperparameter optimization
- ⑦ Results

Table of Contents

- 1 Introduction
- 2 Data Structure
- 3 Algorithms
- 4 Parallelization strategy
- 5 Synchronization strategy
- 6 Hyperparameter optimization
- 7 Results

- **Goal:** test sequential and parallel graph coloring algorithms, analyzing their performances in terms of time, number of colors used and memory occupation
- **Language:** C++
- **HW support:** i7-8550U, 1.80 GHz, 16 GB RAM
- **Operating System:** Windows
- **Reference:** Allwright et al., 1995

Table of Contents

- ① Introduction
- ② Data Structure
- ③ Algorithms
- ④ Parallelization strategy
- ⑤ Synchronization strategy
- ⑥ Hyperparameter optimization
- ⑦ Results

```
vector<int> _colors, _weights, _tmp_degree, _new_colors, _new_weights;  
  
/* one single array containing adjacencies for each node (one vector for each node) */  
vector<vector<int>> _edges;
```

Figure: Data structure from class *graph*

- adjacency list with vectors (one vector with adjacencies for each node)
- additional vectors for colors, weights and other information necessary for the algorithms

Table of Contents

- ① Introduction
- ② Data Structure
- ③ Algorithms**
- ④ Parallelization strategy
- ⑤ Synchronization strategy
- ⑥ Hyperparameter optimization
- ⑦ Results

Greedy

Algorithm 1 Greedy algorithm

```
Given  $G = (V, E)$ 
 $n = |V|$ 
choose a random permutation of the vertices  $V$ 
for  $i=1$  to  $n$  do
    select  $v_i$ 
     $C =$  colors of all colored neighbors of  $v_i$ 
     $c =$  smallest color not in  $C$ 
    color  $v_i$  with color  $c$ 
     $U = U - v_i$ 
end for
```

- sequential
- fast
- usually bad solutions

Jones-Plassman

Algorithm 2 Jones-Plassman

```
assign random weights to each node  $w$ 
 $U := V$ 
while  $|U| > 0$  do
    for all vertices  $v \in U$  do in parallel
         $I := \{v \text{ such that } w(v) > w(u) \text{ for all neighbors } u \in U\}$ 
        for all vertices  $v' \in I$  do in parallel
             $S := \{\text{colors of all neighbors of } v'\}$ 
             $c(v') :=$  minimum color not in  $S$ 
        end for
    end for
     $U = U - I$ 
end while
```

- parallel
- fast
- usually bad solutions

Largest Degree First

- same as JP
- weights are degrees of nodes
- better solutions than JP

Smallest Degree Last

- more complex weight assignment
- better solutions than JP and LDF
- slower

Algorithm 3 Assignment of weights in SDL

```
 $k = 1$   
 $i = 1$   
 $U := V$   
while  $|U| > 0$  do  
  while  $\exists$  vertices  $v \in U$  with  $d^U(v) \leq k$  do  
     $S := \{\text{all vertices } v \text{ with } d^U(v) \leq k\}$   
    for all vertices  $v \in S$  do  
       $w(v) = i$   
    end for  
     $U := U - S$   
     $i := i + 1$   
  end while  
   $k = k + 1$   
end while
```

Table of Contents

- 1 Introduction
- 2 Data Structure
- 3 Algorithms
- 4 Parallelization strategy**
- 5 Synchronization strategy
- 6 Hyperparameter optimization
- 7 Results

Parallelization strategy

- Each thread is assigned a group of nodes
- Coloring is splitted in two distinct phases:
 - find nodes to color (check whether they are local maxima)
 - color the nodes (assign the minimum available color)
- Three different implementations:
 - **No threadpool:** one main thread creates worker threads that either find nodes to color or color them, and collects them after each iteration.
 - **Threadpool:** worker threads are created at the beginning of the algorithm. The main thread schedules jobs that are executed by workers. When the coloring is completed, the threads terminate.
 - **Find and color:** each job determines which nodes can be colored, stores them in a local queue and then colors them.

Table of Contents

- ① Introduction
- ② Data Structure
- ③ Algorithms
- ④ Parallelization strategy
- ⑤ Synchronization strategy**
- ⑥ Hyperparameter optimization
- ⑦ Results

Three cases:



- **No threadpool:** condition variable with counter of active threads to keep the number of active threads constant
- **Threadpool:**
 - condition variable to manage the queue of jobs
 - condition variable with counter of scheduled jobs to wait for their termination
- **Find and color:** condition variable with counter of active jobs to synchronize threads after finding nodes to color (like a barrier). The same condition variable is used, with a different counter, by the main thread to wait for the termination of all jobs

Table of Contents

- ① Introduction
- ② Data Structure
- ③ Algorithms
- ④ Parallelization strategy
- ⑤ Synchronization strategy
- ⑥ Hyperparameter optimization**
- ⑦ Results

Hyperparameter optimization - 1

Number of threads





- best time for 8 threads
- better coloring for less than 8 threads
- final choice: 8 threads (equal to the hardware concurrency)

Hyperparameter optimization - 2

Number of nodes per thread

$$nodes_per_thread = \frac{nodes_in_the_graph}{coef * number_of_threads}$$





- not so relevant
- final choice: $coef = 10$

Table of Contents



- ① Introduction
- ② Data Structure
- ③ Algorithms
- ④ Parallelization strategy
- ⑤ Synchronization strategy
- ⑥ Hyperparameter optimization
- ⑦ Results**

Results - 1





- LDF and SDL find better coloring (less colors) than JP and Greedy
- general exponential correlation between number of colors and number of nodes
- time advantages of parallelism become more evident with big graphs

Results - 2



- parallel implementations of the same algorithms are faster, especially for big graphs
- standard and *find and color* implementations are similar. Standard is better for bigger graphs, *find and color* for smaller ones

Results - 3



- Using a threadpool provides a relevant advantage, especially for small graphs
- Memory occupancy grows significantly for big graphs